
Data Science - A practical Approach

Lorenz Feyen

Sep 24, 2021

CONTENTS

| | | |
|------------|---------------------------------|-----------|
| I | 1. Introduction | 3 |
| 1 | Introduction | 5 |
| II | 2. Data Preparation | 7 |
| 2 | Data Preparation | 9 |
| 3 | Indexing and slicing | 11 |
| 4 | Missing Data | 15 |
| 4.1 | Kamyr digester | 15 |
| 4.2 | Travel times | 17 |
| 4.3 | Material properties | 22 |
| 5 | Concatenation and deduplication | 25 |
| 6 | Some practice | 27 |
| III | 3. Data Preprocessing | 29 |
| 7 | Data Preprocessing | 31 |
| IV | 4. Data Exploration | 33 |
| 8 | Data Exploration | 35 |
| V | 5. Data Visualisation | 37 |
| 9 | Data Visualisation | 39 |
| VI | 6. Machine Learning | 41 |
| 10 | Machine Learning | 43 |

this is a foreword

pdf version can be found here [here](#).

Part I

1. Introduction

INTRODUCTION

this is an introduction

Part II

2. Data Preparation

DATA PREPARATION

this is an introduction

INDEXING AND SLICING

```
import pandas as pd
```

```
min_temp_df = pd.read_csv('https://raw.githubusercontent.com/jbrownlee/Datasets/  
master/daily-min-temperatures.csv')  
min_temp_df
```

```
      Date  Temp  
0  1981-01-01  20.7  
1  1981-01-02  17.9  
2  1981-01-03  18.8  
3  1981-01-04  14.6  
4  1981-01-05  15.8  
...      ...   ...  
3645 1990-12-27  14.0  
3646 1990-12-28  13.6  
3647 1990-12-29  13.5  
3648 1990-12-30  15.7  
3649 1990-12-31  13.0
```

```
[3650 rows x 2 columns]
```

```
min_temp_df.Date = pd.to_datetime(min_temp_df.Date)
```

```
min_temp_df = min_temp_df.set_index('Date')
```

```
min_temp_df.loc['1989-06-01':'1989-06-30']
```

```
      Temp  
Date  
1989-06-01  2.3  
1989-06-02  1.4  
1989-06-03  2.1  
1989-06-04  6.6  
1989-06-05  8.9  
1989-06-06  7.8  
1989-06-07  9.0  
1989-06-08 10.3  
1989-06-09  7.9  
1989-06-10  7.2  
1989-06-11  8.6  
1989-06-12  8.8
```

(continues on next page)

(continued from previous page)

```

1989-06-13    6.2
1989-06-14    9.5
1989-06-15   10.2
1989-06-16    9.7
1989-06-17   11.2
1989-06-18   10.2
1989-06-19   10.1
1989-06-20    8.1
1989-06-21    6.6
1989-06-22    5.0
1989-06-23    4.7
1989-06-24    5.3
1989-06-25    4.5
1989-06-26    2.3
1989-06-27    1.4
1989-06-28    0.5
1989-06-29    2.4
1989-06-30    8.0

```

```
min_temp_df.loc['1989-06-01':'1989-06-30'].mean()
```

```

Temp      6.56
dtype: float64

```

```
import seaborn as sns
```

```

tip_df = sns.load_dataset('tips')
tip_df.head()

```

```

   total_bill  tip  sex smoker  day  time  size
0      16.99  1.01 Female    No  Sun  Dinner    2
1      10.34  1.66   Male    No  Sun  Dinner    3
2      21.01  3.50   Male    No  Sun  Dinner    3
3      23.68  3.31   Male    No  Sun  Dinner    2
4      24.59  3.61 Female    No  Sun  Dinner    4

```

```
tip_index_df = tip_df.set_index('day')
```

```
tip_index_df.loc['Sun']
```

```

   total_bill  tip  sex smoker  time  size
day
Sun      16.99  1.01 Female    No  Dinner    2
Sun      10.34  1.66   Male    No  Dinner    3
Sun      21.01  3.50   Male    No  Dinner    3
Sun      23.68  3.31   Male    No  Dinner    2
Sun      24.59  3.61 Female    No  Dinner    4
..         ...   ...   ...   ...   ...   ...
Sun      20.90  3.50 Female   Yes  Dinner    3
Sun      30.46  2.00   Male   Yes  Dinner    5
Sun      18.15  3.50 Female   Yes  Dinner    3
Sun      23.10  4.00   Male   Yes  Dinner    3
Sun      15.69  1.50   Male   Yes  Dinner    2

```

(continues on next page)

(continued from previous page)

```
[76 rows x 6 columns]
```

```
tip_index_df = tip_df.set_index(['day', 'time'])
```

```
tip_index_df.loc[('Thur', 'Lunch')].tip.mean()
```

```
/tmp/ipykernel_36874/2537502835.py:1: PerformanceWarning: indexing past lexsort depth
↳may impact performance.
  tip_index_df.loc[('Thur', 'Lunch')].tip.mean()
```

```
2.767704918032786
```

```
pd.pivot_table(tip_df, values='total_bill', index='day', columns='time', aggfunc=
↳'median')
```

| time | Lunch | Dinner |
|------|-------|--------|
| day | | |
| Thur | 16.00 | 18.780 |
| Fri | 13.42 | 18.665 |
| Sat | NaN | 18.240 |
| Sun | NaN | 19.630 |

MISSING DATA

In this notebook we will look at a few datasets where values from columns are missing. It is crucial for data science and machine learning to have a dataset where no values are missing as algorithms are usually not able to handle data with information missing.

For python, we will be using the pandas library to handle our dataset.

```
import pandas as pd
```

4.1 Kamyr digester

The first dataset we will be looking at is taken from a physical device equipped with numerous sensors, each timepoint (1 hour) these sensors are read out and the data is collected. Let's have a look at the general structure

```
kamyr_df = pd.read_csv('https://openmv.net/file/kamyr-digester.csv')
kamyr_df.head()
```

| Observation | Y-Kappa | ChipRate | BF-CMratio | BlowFlow | ChipLevel4 | \ | |
|-------------|---------------|---------------|---------------|---------------|-------------|-------------|---|
| 0 | 31-00:00 | 23.10 | 16.520 | 121.717 | 1177.607 | 169.805 | |
| 1 | 31-01:00 | 27.60 | 16.810 | 79.022 | 1328.360 | 341.327 | |
| 2 | 31-02:00 | 23.19 | 16.709 | 79.562 | 1329.407 | 239.161 | |
| 3 | 31-03:00 | 23.60 | 16.478 | 81.011 | 1334.877 | 213.527 | |
| 4 | 31-04:00 | 22.90 | 15.618 | 93.244 | 1334.168 | 243.131 | |
| | | | | | | | |
| | T-upperExt-2 | T-lowerExt-2 | UCZAA | WhiteFlow-4 | ... | SteamFlow-4 | \ |
| 0 | 358.282 | 329.545 | 1.443 | 599.253 | ... | 67.122 | |
| 1 | 351.050 | 329.067 | 1.549 | 537.201 | ... | 60.012 | |
| 2 | 350.022 | 329.260 | 1.600 | 549.611 | ... | 61.304 | |
| 3 | 350.938 | 331.142 | 1.604 | 623.362 | ... | 68.496 | |
| 4 | 351.640 | 332.709 | NaN | 638.672 | ... | 70.022 | |
| | | | | | | | |
| | Lower-HeatT-3 | Upper-HeatT-3 | ChipMass-4 | WeakLiquorF | BlackFlow-2 | \ | |
| 0 | 329.432 | 303.099 | 175.964 | 1127.197 | 1319.039 | | |
| 1 | 330.823 | 304.879 | 163.202 | 665.975 | 1297.317 | | |
| 2 | 329.140 | 303.383 | 164.013 | 677.534 | 1327.072 | | |
| 3 | 328.875 | 302.254 | 181.487 | 767.853 | 1324.461 | | |
| 4 | 328.352 | 300.954 | 183.929 | 888.448 | 1343.424 | | |
| | | | | | | | |
| | WeakWashF | SteamHeatF-3 | T-Top-Chips-4 | SulphidityL-4 | | | |
| 0 | 257.325 | 54.612 | 252.077 | NaN | | | |
| 1 | 241.182 | 46.603 | 251.406 | 29.11 | | | |
| 2 | 237.272 | 51.795 | 251.335 | NaN | | | |

(continues on next page)

(continued from previous page)

```
3      239.478      54.846      250.312      29.02
4      215.372      54.186      249.916      29.01

[5 rows x 23 columns]
```

Interesting, there seem to be 22 sensor values and 1 timestamp for each record. As mechanical devices are prone to noise and dropouts of sensors we would be foolish to assume no missing values are present.

```
kamyr_df.isna().sum().divide(len(kamyr_df)).round(4)*100
```

```
Observation      0.00
Y-Kappa          0.00
ChipRate         1.33
BF-CMratio       4.65
BlowFlow         4.32
ChipLevel4       0.33
T-upperExt-2     0.33
T-lowerExt-2     0.33
UCZAA           7.97
WhiteFlow-4      0.33
AAWhiteSt-4     46.84
AA-Wood-4        0.33
ChipMoisture-4   0.33
SteamFlow-4      0.33
Lower-HeatT-3    0.33
Upper-HeatT-3    0.33
ChipMass-4       0.33
WeakLiquorF      0.33
BlackFlow-2      0.33
WeakWashF        0.33
SteamHeatF-3     0.33
T-Top-Chips-4    0.33
SulphidityL-4    46.84
dtype: float64
```

As expected, the datapoint 'AAWhiteSt-4' even has 46% of data missing! It seems we only have 300 datapoints and presumably these missing values occur in different records our dataset will be decimated if we just drop all rows with missing values.

```
kamyr_df.shape
```

```
(301, 23)
```

```
kamyr_df.dropna().shape
```

```
(131, 23)
```

As we drop all rows with missing values, we are left with only 131 records. Whilst this might be good enough for some purposes, there are more viable options.

Perhaps we can first remove the column with the most missing values and then drop all remaining

```
kamyr_df.drop(columns=['AAWhiteSt-4 ', 'SulphidityL-4 ']).dropna().shape
```

```
(263, 21)
```

Significantly better, although we lost the information of 2 sensors we now have a complete dataset with 263 records. For purposes where those 2 sensors are irrelevant this is a viable option, keep in mind that this dataset is still 100% truthful, as we have not imputed any values.

Another option, where we retain all our records would be using the timely nature of our dataset, each record is a measurement with an interval of 1 hour. I have no knowledge of this dataset but one might make the assumption that the interval of 1 hour is taken as the state of the machine does not alter much in 1 hour. Therefore we could do what is called a forward fill, where we fill in the missing values with the same value of the sensor for the previous measurement.

This would solve nearly all nan values as there might be a problem where the first value is missing. This is shown below.

```
kamyr_df.fillna(method='ffill')['SulphidityL-4 ']
```

```
0      NaN
1    29.11
2    29.11
3    29.02
4    29.01
...
296   30.43
297   30.29
298   30.47
299   30.47
300   30.46
Name: SulphidityL-4 , Length: 301, dtype: float64
```

Although our dataset is not fully the truth, we can see that little to no changes occur in the sensor and using a forward fill is arguably the most suitable option.

4.2 Travel times

Another dataset from the same source contains a collection of recorded travel times and specific information about the travel itself as e.g.: the day of the week, where they were going, ...

```
travel_df = pd.read_csv('https://openmv.net/file/travel-times.csv')
travel_df
```

| | Date | StartTime | DayOfWeek | GoingTo | Distance | MaxSpeed | AvgSpeed | \ |
|-----|----------------|-------------|-----------|------------|------------|----------|----------|---|
| 0 | 1/6/2012 | 16:37 | Friday | Home | 51.29 | 127.4 | 78.3 | |
| 1 | 1/6/2012 | 08:20 | Friday | GSK | 51.63 | 130.3 | 81.8 | |
| 2 | 1/4/2012 | 16:17 | Wednesday | Home | 51.27 | 127.4 | 82.0 | |
| 3 | 1/4/2012 | 07:53 | Wednesday | GSK | 49.17 | 132.3 | 74.2 | |
| 4 | 1/3/2012 | 18:57 | Tuesday | Home | 51.15 | 136.2 | 83.4 | |
| .. | ... | ... | ... | ... | ... | ... | ... | |
| 200 | 7/18/2011 | 08:09 | Monday | GSK | 54.52 | 125.6 | 49.9 | |
| 201 | 7/14/2011 | 08:03 | Thursday | GSK | 50.90 | 123.7 | 76.2 | |
| 202 | 7/13/2011 | 17:08 | Wednesday | Home | 51.96 | 132.6 | 57.5 | |
| 203 | 7/12/2011 | 17:51 | Tuesday | Home | 53.28 | 125.8 | 61.6 | |
| 204 | 7/11/2011 | 16:56 | Monday | Home | 51.73 | 125.0 | 62.8 | |
| | AvgMovingSpeed | FuelEconomy | TotalTime | MovingTime | Take407All | Comments | | |
| 0 | 84.8 | NaN | 39.3 | 36.3 | No | NaN | | |

(continues on next page)

(continued from previous page)

| | | | | | | |
|-------------------------|------|------|------|------|-----|-----|
| 1 | 88.9 | NaN | 37.9 | 34.9 | No | NaN |
| 2 | 85.8 | NaN | 37.5 | 35.9 | No | NaN |
| 3 | 82.9 | NaN | 39.8 | 35.6 | No | NaN |
| 4 | 88.1 | NaN | 36.8 | 34.8 | No | NaN |
| .. | ... | ... | ... | ... | ... | ... |
| 200 | 82.4 | 7.89 | 65.5 | 39.7 | No | NaN |
| 201 | 95.1 | 7.89 | 40.1 | 32.1 | Yes | NaN |
| 202 | 76.7 | NaN | 54.2 | 40.6 | Yes | NaN |
| 203 | 87.6 | NaN | 51.9 | 36.5 | Yes | NaN |
| 204 | 92.5 | NaN | 49.5 | 33.6 | Yes | NaN |
| [205 rows x 13 columns] | | | | | | |

we have a total of 205 records and we can already see that the FuelEconomy column seems pretty bad, let's quantify that.

```
travel_df.isna().sum().divide(len(travel_df)).round(4)*100
```

```
Date          0.00
StartTime      0.00
DayOfWeek      0.00
GoingTo        0.00
Distance       0.00
MaxSpeed       0.00
AvgSpeed       0.00
AvgMovingSpeed 0.00
FuelEconomy    8.29
TotalTime      0.00
MovingTime     0.00
Take407All     0.00
Comments      88.29
dtype: float64
```

In the end, it doesn't seem that bad, but there are comments and nearly none of them are filled in. Which in perspective is understandable. Let's see what the comments look like

```
travel_df[~travel_df.Comments.isna()].Comments
```

```
15          Put snow tires on
39          Heavy rain
49          Huge traffic backup
50    Pumped tires up: check fuel economy improved?
52          Backed up at Bronte
54          Backed up at Bronte
60          Rainy
78          Rain, rain, rain
91          Rain, rain, rain
92    Accident: backup from Hamilton to 407 ramp
110         Raining
132         Back to school traffic?
133    Took 407 all the way (to McMaster)
150         Heavy volume on Derry
156         Start early to run a batch
158    Accident at 403/highway 6; detour along Dundas
165         Detour taken
166         Must be Friday
172         Medium amount of rain
```

(continues on next page)

(continued from previous page)

```

174                                New tires
182                        Turn around on Derry
184                        Empty roads
187                        Police slowdown on 403
189                        Accident blocked 407 exit
Name: Comments, dtype: object

```

As you would expect, these comments are text based. Now imagine we would like to run some Natural Language Processing (NLP) on these, it would be a pain to perform string operations on it when it is riddled with missing values.

Here a simple example where we select all records containing the word 'rain', with no avail.

```
travel_df[travel_df.Comments.str.lower().str.contains('rain')]
```

```

-----
ValueError                                Traceback (most recent call last)
/tmp/ipykernel_36906/1298831137.py in <module>
----> 1 travel_df[travel_df.Comments.str.lower().str.contains('rain')]

~/git/data-science-practical-approach/venv/lib/python3.8/site-packages/pandas/core/
frame.py in __getitem__(self, key)
   3446
   3447     # Do we have a (boolean) 1d indexer?
-> 3448     if com.is_bool_indexer(key):
   3449         return self._getitem_bool_array(key)
   3450

~/git/data-science-practical-approach/venv/lib/python3.8/site-packages/pandas/core/
common.py in is_bool_indexer(key)
   137         # Don't raise on e.g. ["A", "B", np.nan], see
   138         # test_loc_getitem_list_of_labels_categoricalindex_with_
-> na
   139         raise ValueError(na_msg)
   140         return False
   141         return True

ValueError: Cannot mask with non-boolean array containing NA / NaN values

```

The last line of the python error traceback gives us the reason it failed, because there were NaN values present.

Luckily the string variable has more or less it's on 'null' value, being an empty string, this way these operations are still possible, most of the comments will just contain nothing.

```
travel_df.Comments = travel_df.Comments.fillna('')
```

```
travel_df[travel_df.Comments.str.lower().str.contains('rain')]
```

| | Date | StartTime | DayOfWeek | GoingTo | Distance | MaxSpeed | AvgSpeed | \ |
|-----|----------------|-------------|-----------|------------|------------|----------|----------|---|
| 39 | 11/29/2011 | 07:23 | Tuesday | GSK | 51.74 | 112.2 | 55.3 | |
| 60 | 11/9/2011 | 16:15 | Wednesday | Home | 51.28 | 121.4 | 65.9 | |
| 78 | 10/25/2011 | 17:24 | Tuesday | Home | 52.87 | 123.5 | 65.1 | |
| 91 | 10/12/2011 | 17:47 | Wednesday | Home | 51.40 | 114.4 | 59.7 | |
| 110 | 9/27/2011 | 07:36 | Tuesday | GSK | 50.65 | 128.1 | 86.3 | |
| 172 | 8/9/2011 | 08:15 | Tuesday | GSK | 49.08 | 134.8 | 60.5 | |
| | AvgMovingSpeed | FuelEconomy | TotalTime | MovingTime | Take407All | \ | | |

(continues on next page)

(continued from previous page)

| | | | | | |
|----------|-----------------------|------|------|------|-----|
| 39 | 61.0 | NaN | 56.2 | 50.9 | No |
| 60 | 71.8 | 9.35 | 46.7 | 42.1 | No |
| 78 | 72.4 | 8.97 | 48.7 | 43.8 | No |
| 91 | 65.8 | 8.75 | 51.7 | 46.9 | No |
| 110 | 88.6 | 8.31 | 35.2 | 34.3 | Yes |
| 172 | 67.2 | 8.54 | 48.7 | 43.8 | No |
| Comments | | | | | |
| 39 | Heavy rain | | | | |
| 60 | Rainy | | | | |
| 78 | Rain, rain, rain | | | | |
| 91 | Rain, rain, rain | | | | |
| 110 | Raining | | | | |
| 172 | Medium amount of rain | | | | |

Fixed! now we can use the comments for analysis.

We still have to fix the FuelEconomy, let us take a look at the non NaN values

```
travel_df[~travel_df.FuelEconomy.isna()]
```

| | Date | StartTime | DayOfWeek | GoingTo | Distance | MaxSpeed | AvgSpeed | \ |
|-----|----------------|-------------|-----------|------------|------------|----------|----------|---|
| 6 | 1/2/2012 | 17:31 | Monday | Home | 51.37 | 123.2 | 82.9 | |
| 7 | 1/2/2012 | 07:34 | Monday | GSK | 49.01 | 128.3 | 77.5 | |
| 8 | 12/23/2011 | 08:01 | Friday | GSK | 52.91 | 130.3 | 80.9 | |
| 9 | 12/22/2011 | 17:19 | Thursday | Home | 51.17 | 122.3 | 70.6 | |
| 10 | 12/22/2011 | 08:16 | Thursday | GSK | 49.15 | 129.4 | 74.0 | |
| .. | ... | ... | ... | ... | ... | ... | ... | |
| 197 | 7/20/2011 | 08:24 | Wednesday | GSK | 48.50 | 125.8 | 75.7 | |
| 198 | 7/19/2011 | 17:17 | Tuesday | Home | 51.16 | 126.7 | 92.2 | |
| 199 | 7/19/2011 | 08:11 | Tuesday | GSK | 50.96 | 124.3 | 82.3 | |
| 200 | 7/18/2011 | 08:09 | Monday | GSK | 54.52 | 125.6 | 49.9 | |
| 201 | 7/14/2011 | 08:03 | Thursday | GSK | 50.90 | 123.7 | 76.2 | |
| | | | | | | | | |
| | AvgMovingSpeed | FuelEconomy | TotalTime | MovingTime | Take407All | Comments | | |
| 6 | 87.3 | - | 37.2 | 35.3 | No | | | |
| 7 | 85.9 | - | 37.9 | 34.3 | No | | | |
| 8 | 88.3 | 8.89 | 39.3 | 36.0 | No | | | |
| 9 | 78.1 | 8.89 | 43.5 | 39.3 | No | | | |
| 10 | 81.4 | 8.89 | 39.8 | 36.2 | No | | | |
| .. | ... | ... | ... | ... | ... | ... | | |
| 197 | 87.3 | 7.89 | 38.5 | 33.3 | Yes | | | |
| 198 | 102.6 | 7.89 | 33.3 | 29.9 | Yes | | | |
| 199 | 96.4 | 7.89 | 37.2 | 31.7 | Yes | | | |
| 200 | 82.4 | 7.89 | 65.5 | 39.7 | No | | | |
| 201 | 95.1 | 7.89 | 40.1 | 32.1 | Yes | | | |

[188 rows x 13 columns]

It seems that aside NaN values there are also other intruders, a quick check on the data type (Dtype) reveals it is not recognised as a number!

```
travel_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
```

(continues on next page)

(continued from previous page)

```
Data columns (total 13 columns):
#      Column      Non-Null Count  Dtype
---  -
0     Date         205 non-null    object
1     StartTime     205 non-null    object
2     DayOfWeek      205 non-null    object
3     GoingTo        205 non-null    object
4     Distance       205 non-null    float64
5     MaxSpeed       205 non-null    float64
6     AvgSpeed       205 non-null    float64
7     AvgMovingSpeed  205 non-null    float64
8     FuelEconomy     188 non-null    object
9     TotalTime      205 non-null    float64
10    MovingTime     205 non-null    float64
11    Take407All     205 non-null    object
12    Comments       205 non-null    object
dtypes: float64(6), object(7)
memory usage: 20.9+ KB
```

The column is noted as an object or string type, meaning that these numbers are given as '9.24' instead of 9.24 and numerical operations are not possible. We can cast them to numeric but have to warn pandas to coerce errors, meaning errors will be converted to NaN values. Later we'll handle the NaN's.

```
travel_df.FuelEconomy = pd.to_numeric(travel_df.FuelEconomy, errors='coerce')
travel_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 13 columns):
#      Column      Non-Null Count  Dtype
---  -
0     Date         205 non-null    object
1     StartTime     205 non-null    object
2     DayOfWeek      205 non-null    object
3     GoingTo        205 non-null    object
4     Distance       205 non-null    float64
5     MaxSpeed       205 non-null    float64
6     AvgSpeed       205 non-null    float64
7     AvgMovingSpeed  205 non-null    float64
8     FuelEconomy     186 non-null    float64
9     TotalTime      205 non-null    float64
10    MovingTime     205 non-null    float64
11    Take407All     205 non-null    object
12    Comments       205 non-null    object
dtypes: float64(7), object(6)
memory usage: 20.9+ KB
```

Wonderful, now the column is numerical and we can see 2 more missing values have popped up! We could easily drop these 19 records and have a complete dataset.

```
travel_df.dropna()
```

| | Date | StartTime | DayOfWeek | GoingTo | Distance | MaxSpeed | AvgSpeed | \ |
|----|------------|-----------|-----------|---------|----------|----------|----------|---|
| 8 | 12/23/2011 | 08:01 | Friday | GSK | 52.91 | 130.3 | 80.9 | |
| 9 | 12/22/2011 | 17:19 | Thursday | Home | 51.17 | 122.3 | 70.6 | |
| 10 | 12/22/2011 | 08:16 | Thursday | GSK | 49.15 | 129.4 | 74.0 | |

(continues on next page)

(continued from previous page)

| | | | | | | | |
|-----|----------------|-------------|-----------|------------|------------|----------|------|
| 11 | 12/21/2011 | 07:45 | Wednesday | GSK | 51.77 | 124.8 | 71.7 |
| 12 | 12/20/2011 | 16:05 | Tuesday | Home | 51.45 | 130.1 | 75.2 |
| .. | ... | ... | ... | ... | ... | ... | ... |
| 197 | 7/20/2011 | 08:24 | Wednesday | GSK | 48.50 | 125.8 | 75.7 |
| 198 | 7/19/2011 | 17:17 | Tuesday | Home | 51.16 | 126.7 | 92.2 |
| 199 | 7/19/2011 | 08:11 | Tuesday | GSK | 50.96 | 124.3 | 82.3 |
| 200 | 7/18/2011 | 08:09 | Monday | GSK | 54.52 | 125.6 | 49.9 |
| 201 | 7/14/2011 | 08:03 | Thursday | GSK | 50.90 | 123.7 | 76.2 |
| | | | | | | | |
| | AvgMovingSpeed | FuelEconomy | TotalTime | MovingTime | Take407All | Comments | |
| 8 | 88.3 | 8.89 | 39.3 | 36.0 | No | | |
| 9 | 78.1 | 8.89 | 43.5 | 39.3 | No | | |
| 10 | 81.4 | 8.89 | 39.8 | 36.2 | No | | |
| 11 | 78.9 | 8.89 | 43.3 | 39.4 | No | | |
| 12 | 82.7 | 8.89 | 41.1 | 37.3 | No | | |
| .. | ... | ... | ... | ... | ... | ... | ... |
| 197 | 87.3 | 7.89 | 38.5 | 33.3 | Yes | | |
| 198 | 102.6 | 7.89 | 33.3 | 29.9 | Yes | | |
| 199 | 96.4 | 7.89 | 37.2 | 31.7 | Yes | | |
| 200 | 82.4 | 7.89 | 65.5 | 39.7 | No | | |
| 201 | 95.1 | 7.89 | 40.1 | 32.1 | Yes | | |

[186 rows x 13 columns]

However im leaving them as an exercise for you to apply a technique we will see in the next part

4.3 Material properties

Another dataset from the same source contains the material properties from 30 samples, this time there is not timestamp as the samples are not related in time with each other.

```
material_df = pd.read_csv('http://openmv.net/file/raw-material-properties.csv')
material_df
```

| | Sample | size1 | size2 | size3 | density1 | density2 | density3 |
|----|--------|-------|-------|-------|----------|----------|----------|
| 0 | X12558 | 0.696 | 2.69 | 6.38 | 41.8 | 17.18 | 3.90 |
| 1 | X14728 | 0.636 | 2.30 | 5.14 | 38.1 | 12.73 | 3.89 |
| 2 | X15468 | 0.841 | 2.85 | 5.20 | 37.6 | 13.58 | 3.98 |
| 3 | X21364 | 0.609 | 2.13 | 4.62 | 34.2 | 11.12 | 4.02 |
| 4 | X23671 | 0.684 | 2.16 | 4.87 | 36.4 | 12.24 | 3.92 |
| 5 | X24055 | 0.762 | 2.81 | 6.36 | 38.1 | 13.28 | 3.89 |
| 6 | X24905 | 0.552 | 2.34 | 5.03 | 41.3 | 16.71 | 3.86 |
| 7 | X25917 | 0.501 | 2.17 | 5.09 | NaN | NaN | NaN |
| 8 | X27871 | 0.619 | 2.11 | 5.13 | NaN | NaN | NaN |
| 9 | X28690 | 0.610 | 2.10 | 4.18 | 35.0 | 12.15 | 3.86 |
| 10 | X31385 | 0.532 | 2.09 | 4.93 | NaN | NaN | NaN |
| 11 | X31813 | 0.738 | 2.29 | 5.47 | NaN | NaN | NaN |
| 12 | X32807 | 0.779 | 2.62 | 5.59 | NaN | NaN | NaN |
| 13 | X33943 | 0.537 | 2.23 | 5.41 | 35.2 | 11.34 | 3.99 |
| 14 | X35035 | 0.702 | 2.05 | 5.10 | 34.2 | 10.54 | 4.02 |
| 15 | X39223 | 0.768 | 2.51 | 5.09 | 34.9 | 12.55 | 3.90 |
| 16 | X40503 | 0.714 | 2.56 | 6.03 | 35.6 | 12.20 | 4.02 |
| 17 | X41400 | 0.621 | 2.42 | 5.10 | 38.7 | 14.27 | 3.98 |
| 18 | X42988 | 0.726 | 2.11 | 4.69 | 37.1 | 13.14 | 3.98 |

(continues on next page)

(continued from previous page)

| | | | | | | | |
|----|--------|-------|------|------|------|-------|------|
| 19 | X44749 | 0.698 | 2.36 | 5.40 | 36.6 | 12.16 | 4.01 |
| 20 | X45295 | NaN | NaN | NaN | 38.1 | 13.34 | 3.89 |
| 21 | X46965 | 0.759 | 2.47 | 4.83 | 38.7 | 14.83 | 3.89 |
| 22 | X49666 | 0.535 | 2.13 | 5.23 | NaN | NaN | NaN |
| 23 | X50678 | 0.716 | 2.29 | 5.45 | 37.3 | 13.70 | 3.92 |
| 24 | X52894 | 0.635 | 2.08 | 4.94 | NaN | NaN | NaN |
| 25 | X53925 | 0.598 | 2.12 | 4.69 | 37.9 | 13.45 | 3.78 |
| 26 | X54254 | 0.700 | 2.47 | 5.22 | 38.8 | 14.72 | 3.92 |
| 27 | X54272 | 0.957 | 2.96 | 7.37 | 36.2 | 13.38 | 4.20 |
| 28 | X54394 | 0.759 | 2.66 | 5.36 | 35.2 | 12.19 | 3.98 |
| 29 | X55408 | 0.661 | 2.10 | 4.27 | NaN | NaN | NaN |
| 30 | X56952 | 0.646 | 2.38 | 4.51 | 40.1 | 15.68 | 3.86 |
| 31 | X57095 | 0.662 | 2.34 | 4.71 | 35.0 | 12.37 | 3.90 |
| 32 | X57128 | 0.749 | 2.43 | 5.16 | 37.3 | 13.04 | 3.92 |
| 33 | X61870 | 0.598 | 2.21 | 4.90 | NaN | NaN | NaN |
| 34 | X61888 | 0.619 | 2.59 | 5.81 | NaN | NaN | NaN |
| 35 | X72736 | 0.693 | 2.05 | 5.02 | 39.6 | 15.55 | 3.94 |

let us quantify the amount of missing data

```
material_df.isna().sum().divide(len(material_df)).round(4)*100
```

```
Sample      0.00
size1       2.78
size2       2.78
size3       2.78
density1    27.78
density2    27.78
density3    27.78
dtype: float64
```

Unfortunately that is a lot of missing data, covered in all records, dropping here seems almost impossible if we want to keep a healthy amount of records.

Here it would be wise to go for a more elaborate method of imputation, I opted for the K-nearest neighbours method, which looks at the K most similar records in the dataset to make an educated guess on what the missing value could be, this because we can assume that records with similar data are also similar over all the properties (columns).

Im using the sklearn library for this, which has more imputation techniques such as MICE. More info can be found [here](#)

```
from sklearn.impute import KNNImputer
```

im creating an imputer object and specify that i want to use the 5 most similar records and weigh them by distance from the to imputed record, meaning closer neighbours are more important.

```
imputer = KNNImputer(n_neighbors=5, weights="distance")
```

As the imputer only takes numerical values I had to do some pandas magic and drop the first column, which I then added again. The result is a fully filled dataset, you can recognise the new values as they are not rounded.

```
pd.DataFrame(
    imputer.fit_transform(material_df.drop(columns=['Sample'])),
    columns=material_df.columns.drop('Sample')
)
```

| | size1 | size2 | size3 | density1 | density2 | density3 |
|----|----------|----------|----------|-----------|-----------|----------|
| 0 | 0.696000 | 2.690000 | 6.380000 | 41.800000 | 17.180000 | 3.900000 |
| 1 | 0.636000 | 2.300000 | 5.140000 | 38.100000 | 12.730000 | 3.890000 |
| 2 | 0.841000 | 2.850000 | 5.200000 | 37.600000 | 13.580000 | 3.980000 |
| 3 | 0.609000 | 2.130000 | 4.620000 | 34.200000 | 11.120000 | 4.020000 |
| 4 | 0.684000 | 2.160000 | 4.870000 | 36.400000 | 12.240000 | 3.920000 |
| 5 | 0.762000 | 2.810000 | 6.360000 | 38.100000 | 13.280000 | 3.890000 |
| 6 | 0.552000 | 2.340000 | 5.030000 | 41.300000 | 16.710000 | 3.860000 |
| 7 | 0.501000 | 2.170000 | 5.090000 | 38.495282 | 14.029399 | 3.931180 |
| 8 | 0.619000 | 2.110000 | 5.130000 | 37.405275 | 13.157346 | 3.943667 |
| 9 | 0.610000 | 2.100000 | 4.180000 | 35.000000 | 12.150000 | 3.860000 |
| 10 | 0.532000 | 2.090000 | 4.930000 | 37.811132 | 13.646072 | 3.908364 |
| 11 | 0.738000 | 2.290000 | 5.470000 | 37.088833 | 13.255412 | 3.941654 |
| 12 | 0.779000 | 2.620000 | 5.590000 | 36.540567 | 12.889902 | 3.970973 |
| 13 | 0.537000 | 2.230000 | 5.410000 | 35.200000 | 11.340000 | 3.990000 |
| 14 | 0.702000 | 2.050000 | 5.100000 | 34.200000 | 10.540000 | 4.020000 |
| 15 | 0.768000 | 2.510000 | 5.090000 | 34.900000 | 12.550000 | 3.900000 |
| 16 | 0.714000 | 2.560000 | 6.030000 | 35.600000 | 12.200000 | 4.020000 |
| 17 | 0.621000 | 2.420000 | 5.100000 | 38.700000 | 14.270000 | 3.980000 |
| 18 | 0.726000 | 2.110000 | 4.690000 | 37.100000 | 13.140000 | 3.980000 |
| 19 | 0.698000 | 2.360000 | 5.400000 | 36.600000 | 12.160000 | 4.010000 |
| 20 | 0.733097 | 2.653959 | 5.881504 | 38.100000 | 13.340000 | 3.890000 |
| 21 | 0.759000 | 2.470000 | 4.830000 | 38.700000 | 14.830000 | 3.890000 |
| 22 | 0.535000 | 2.130000 | 5.230000 | 37.391815 | 13.089536 | 3.944335 |
| 23 | 0.716000 | 2.290000 | 5.450000 | 37.300000 | 13.700000 | 3.920000 |
| 24 | 0.635000 | 2.080000 | 4.940000 | 37.254724 | 13.206262 | 3.933904 |
| 25 | 0.598000 | 2.120000 | 4.690000 | 37.900000 | 13.450000 | 3.780000 |
| 26 | 0.700000 | 2.470000 | 5.220000 | 38.800000 | 14.720000 | 3.920000 |
| 27 | 0.957000 | 2.960000 | 7.370000 | 36.200000 | 13.380000 | 4.200000 |
| 28 | 0.759000 | 2.660000 | 5.360000 | 35.200000 | 12.190000 | 3.980000 |
| 29 | 0.661000 | 2.100000 | 4.270000 | 36.172345 | 12.755632 | 3.887375 |
| 30 | 0.646000 | 2.380000 | 4.510000 | 40.100000 | 15.680000 | 3.860000 |
| 31 | 0.662000 | 2.340000 | 4.710000 | 35.000000 | 12.370000 | 3.900000 |
| 32 | 0.749000 | 2.430000 | 5.160000 | 37.300000 | 13.040000 | 3.920000 |
| 33 | 0.598000 | 2.210000 | 4.900000 | 37.865882 | 13.826029 | 3.887021 |
| 34 | 0.619000 | 2.590000 | 5.810000 | 35.932339 | 12.318210 | 3.989911 |
| 35 | 0.693000 | 2.050000 | 5.020000 | 39.600000 | 15.550000 | 3.940000 |

This concludes the part of missing values, perhaps you can try yourself and impute the missing values for the FuelEconomy using the SimpleImputer or even the IterativeImputer.

CONCATENATION AND DEDUPLICATION

https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2020-01.csv

```
import pandas as pd
```

```
-----  
ModuleNotFoundError                                Traceback (most recent call last)  
/tmp/ipykernel_13304/4080736814.py in <module>  
----> 1 import pandas as pd  
  
ModuleNotFoundError: No module named 'pandas'
```

```
df = pd.read_csv('https://s3.amazonaws.com/nyc-tlc/trip+data/yellow_tripdata_2020-01.  
↳csv')
```

```
/home/lorenz/.local/lib/python3.8/site-packages/IPython/core/interactiveshell.  
↳py:3441: DtypeWarning: Columns (6) have mixed types.Specify dtype option on import  
↳or set low_memory=False.  
exec(code_obj, self.user_global_ns, self.user_ns)
```

```
df
```

```
VendorID tpep_pickup_datetime tpep_dropoff_datetime passenger_count \  
0          1.0  2020-01-01 00:28:15    2020-01-01 00:33:03          1.0  
1          1.0  2020-01-01 00:35:39    2020-01-01 00:43:04          1.0  
2          1.0  2020-01-01 00:47:41    2020-01-01 00:53:52          1.0  
3          1.0  2020-01-01 00:55:23    2020-01-01 01:00:14          1.0  
4          2.0  2020-01-01 00:01:58    2020-01-01 00:04:16          1.0  
...          ...          ...          ...  
6405003      NaN  2020-01-31 22:51:00    2020-01-31 23:22:00          NaN  
6405004      NaN  2020-01-31 22:10:00    2020-01-31 23:26:00          NaN  
6405005      NaN  2020-01-31 22:50:07    2020-01-31 23:17:57          NaN  
6405006      NaN  2020-01-31 22:25:53    2020-01-31 22:48:32          NaN  
6405007      NaN  2020-01-31 22:44:00    2020-01-31 23:06:00          NaN  
  
trip_distance  RatecodeID store_and_fwd_flag  PULocationID \  
0              1.20          1.0              N           238  
1              1.20          1.0              N           239  
2              0.60          1.0              N           238  
3              0.80          1.0              N           238  
4              0.00          1.0              N           193  
...          ...          ...          ...  
6405003        3.24          NaN            NaN           237
```

(continues on next page)

(continued from previous page)

| | | | | |
|---------|-------|-----|-----|-----|
| 6405004 | 22.13 | NaN | NaN | 259 |
| 6405005 | 10.51 | NaN | NaN | 137 |
| 6405006 | 5.49 | NaN | NaN | 50 |
| 6405007 | 11.60 | NaN | NaN | 179 |

| | DOLocationID | payment_type | fare_amount | extra | mta_tax | tip_amount | \ |
|---------|--------------|--------------|-------------|-------|---------|------------|---|
| 0 | 239 | 1.0 | 6.00 | 3.00 | 0.5 | 1.47 | |
| 1 | 238 | 1.0 | 7.00 | 3.00 | 0.5 | 1.50 | |
| 2 | 238 | 1.0 | 6.00 | 3.00 | 0.5 | 1.00 | |
| 3 | 151 | 1.0 | 5.50 | 0.50 | 0.5 | 1.36 | |
| 4 | 193 | 2.0 | 3.50 | 0.50 | 0.5 | 0.00 | |
| ... | ... | ... | ... | ... | ... | ... | |
| 6405003 | 234 | NaN | 17.59 | 2.75 | 0.5 | 0.00 | |
| 6405004 | 45 | NaN | 46.67 | 2.75 | 0.5 | 0.00 | |
| 6405005 | 169 | NaN | 48.85 | 2.75 | 0.0 | 0.00 | |
| 6405006 | 42 | NaN | 27.17 | 2.75 | 0.0 | 0.00 | |
| 6405007 | 205 | NaN | 54.56 | 2.75 | 0.5 | 0.00 | |

| | tolls_amount | improvement_surcharge | total_amount | \ |
|---------|--------------|-----------------------|--------------|---|
| 0 | 0.00 | 0.3 | 11.27 | |
| 1 | 0.00 | 0.3 | 12.30 | |
| 2 | 0.00 | 0.3 | 10.80 | |
| 3 | 0.00 | 0.3 | 8.16 | |
| 4 | 0.00 | 0.3 | 4.80 | |
| ... | ... | ... | ... | |
| 6405003 | 0.00 | 0.3 | 21.14 | |
| 6405004 | 12.24 | 0.3 | 62.46 | |
| 6405005 | 0.00 | 0.3 | 51.90 | |
| 6405006 | 0.00 | 0.3 | 30.22 | |
| 6405007 | 0.00 | 0.3 | 58.11 | |

| | congestion_surcharge |
|---------|----------------------|
| 0 | 2.5 |
| 1 | 2.5 |
| 2 | 2.5 |
| 3 | 0.0 |
| 4 | 0.0 |
| ... | ... |
| 6405003 | 0.0 |
| 6405004 | 0.0 |
| 6405005 | 0.0 |
| 6405006 | 0.0 |
| 6405007 | 0.0 |

[6405008 rows x 18 columns]

SOME PRACTICE

Now that you have learned techniques in data preparation, why don't you put them to use in this wonderfully horrifying dataset. Good luck!

```
import os
import json

import pandas as pd
import kaggle
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
/tmp/ipykernel_36949/2054829274.py in <module>
      3
      4 import pandas as pd
----> 5 import kaggle

ModuleNotFoundError: No module named 'kaggle'
```

```
if not os.path.exists("/root/.kaggle"):
    os.mkdir("/root/.kaggle")

with open('/root/.kaggle/kaggle.json', 'w') as f:
    json.dump(
        {
            "username": "lorenzof",
            "key": "7a44a9e99b27e796177d793a3d85b8cf"
        }, f)
```

```
kaggle.api.dataset_download_files(dataset='PromptCloudHQ/us-jobs-on-monstercom', path=
↳ './data', unzip=True)
```

```
df = pd.read_csv('./data/monster_com-job_sample.csv')
```

```
df.head()
```

| | country | country_code | date_added | has_expired | \ |
|---|--------------------------|--------------|------------|-------------|---|
| 0 | United States of America | US | NaN | No | |
| 1 | United States of America | US | NaN | No | |
| 2 | United States of America | US | NaN | No | |
| 3 | United States of America | US | NaN | No | |

(continues on next page)

(continued from previous page)

| | | | | |
|---|---|---|---------------------|----|
| 4 | United States of America | US | NaN | No |
| | job_board | | job_description | \ |
| 0 | jobs.monster.com | TeamSoft is seeing an IT Support Specialist to... | | |
| 1 | jobs.monster.com | The Wisconsin State Journal is seeking a flexi... | | |
| 2 | jobs.monster.com | Report this job About the Job DePuy Synthes Co... | | |
| 3 | jobs.monster.com | Why Join Altec? If you're considering a career... | | |
| 4 | jobs.monster.com | Position ID# 76162 # Positions 1 State CT C... | | |
| | job_title | | job_type | \ |
| 0 | IT Support Technician Job in Madison | | Full Time Employee | |
| 1 | Business Reporter/Editor Job in Madison | | Full Time | |
| 2 | Johnson & Johnson Family of Companies Job Appl... | | Full Time, Employee | |
| 3 | Engineer - Quality Job in Dixon | | Full Time | |
| 4 | Shift Supervisor - Part-Time Job in Camphill | | Full Time Employee | |
| | location | | | \ |
| 0 | Madison, WI 53702 | | | |
| 1 | Madison, WI 53708 | | | |
| 2 | DePuy Synthes Companies is a member of Johnson... | | | |
| 3 | Dixon, CA | | | |
| 4 | Camphill, PA | | | |
| | organization | | | \ |
| 0 | NaN | | | |
| 1 | Printing and Publishing | | | |
| 2 | Personal and Household Services | | | |
| 3 | Altec Industries | | | |
| 4 | Retail | | | |
| | page_url | salary | | \ |
| 0 | http://jobview.monster.com/it-support-technici... | NaN | | |
| 1 | http://jobview.monster.com/business-reporter-e... | NaN | | |
| 2 | http://jobview.monster.com/senior-training-lea... | NaN | | |
| 3 | http://jobview.monster.com/engineer-quality-jo... | NaN | | |
| 4 | http://jobview.monster.com/shift-supervisor-pa... | NaN | | |
| | sector | | uniq_id | |
| 0 | IT/Software Development | 11d599f229a80023d2f40e7c52cd941e | | |
| 1 | NaN | e4cbb126dabf22159aff90223243ff2a | | |
| 2 | NaN | 839106b353877fa3d896ffb9c1fe01c0 | | |
| 3 | Experienced (Non-Manager) | 58435fcab804439efdcaa7ecca0fd783 | | |
| 4 | Project/Program Management | 64d0272dc8496abfd9523a8df63c184c | | |

Need some inspiration? perhaps [this](#) might work.

Part III

3. Data Preprocessing

DATA PREPROCESSING

this is an introduction

Part IV

4. Data Exploration

DATA EXPLORATION

this is an introduction

Part V

5. Data Visualisation

DATA VISUALISATION

this is an introduction

Part VI

6. Machine Learning

MACHINE LEARNING

this is an introduction