
Data Science - A practical Approach

Lorenz Feyen

Sep 27, 2021

CONTENTS

I	1. Introduction	3
1	Introduction	5
II	2. Data Preparation	7
2	Data Preparation	9
3	Missing Data	11
3.1	Kamyr digester	11
3.2	Travel times	13
3.3	Material properties	18
4	Concatenation and deduplication	21
4.1	Concatenation	21
4.2	Deduplication	22
5	Outliers and validity	27
6	Some practice	29
III	3. Data Preprocessing	31
7	Data Preprocessing	33
8	Indexing and slicing	35
IV	4. Data Exploration	39
9	Data Exploration	41
V	5. Data Visualisation	43
10	Data Visualisation	45
VI	6. Machine Learning	47
11	Machine Learning	49

this is a foreword

pdf version can be found [here](#).

Part I

1. Introduction

INTRODUCTION

this is an introduction

Part II

2. Data Preparation

DATA PREPARATION

this is an introduction

MISSING DATA

In this notebook we will look at a few datasets where values from columns are missing. It is crucial for data science and machine learning to have a dataset where no values are missing as algorithms are usually not able to handle data with information missing.

For python, we will be using the pandas library to handle our dataset.

```
import pandas as pd
```

3.1 Kamyr digester

The first dataset we will be looking at is taken from a physical device equipped with numerous sensors, each timepoint (1 hour) these sensors are read out and the data is collected. Let's have a look at the general structure

```
kamyr_df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-
practical-approach/main/src/c2_data_preparation/data/kamyr-digester.csv')
kamyr_df.head()
```

Observation	Y-Kappa	ChipRate	BF-CMratio	BlowFlow	ChipLevel4	\
0	31-00:00	23.10	16.520	121.717	1177.607	169.805
1	31-01:00	27.60	16.810	79.022	1328.360	341.327
2	31-02:00	23.19	16.709	79.562	1329.407	239.161
3	31-03:00	23.60	16.478	81.011	1334.877	213.527
4	31-04:00	22.90	15.618	93.244	1334.168	243.131
T-upperExt-2	T-lowerExt-2	UCZAA	WhiteFlow-4	...	SteamFlow-4	\
0	358.282	329.545	1.443	599.253	...	67.122
1	351.050	329.067	1.549	537.201	...	60.012
2	350.022	329.260	1.600	549.611	...	61.304
3	350.938	331.142	1.604	623.362	...	68.496
4	351.640	332.709	NaN	638.672	...	70.022
Lower-HeatT-3	Upper-HeatT-3	ChipMass-4	WeakLiquorF	BlackFlow-2	\	
0	329.432	303.099	175.964	1127.197	1319.039	
1	330.823	304.879	163.202	665.975	1297.317	
2	329.140	303.383	164.013	677.534	1327.072	
3	328.875	302.254	181.487	767.853	1324.461	
4	328.352	300.954	183.929	888.448	1343.424	
WeakWashF	SteamHeatF-3	T-Top-Chips-4	SulphidityL-4			
0	257.325	54.612	252.077	NaN		
1	241.182	46.603	251.406	29.11		

(continues on next page)

(continued from previous page)

2	237.272	51.795	251.335	NaN
3	239.478	54.846	250.312	29.02
4	215.372	54.186	249.916	29.01
[5 rows x 23 columns]				

Interesting, there seem to be 22 sensor values and 1 timestamp for each record. As mechanical devices are prone to noise and dropouts of sensors we would be foolish to assume no missing values are present.

```
kamyr_df.isna().sum().divide(len(kamyr_df)).round(4)*100
```

```
Observation      0.00
Y-Kappa          0.00
ChipRate         1.33
BF-CMratio       4.65
BlowFlow         4.32
ChipLevel4       0.33
T-upperExt-2     0.33
T-lowerExt-2     0.33
UCZAA           7.97
WhiteFlow-4      0.33
AAWhiteSt-4     46.84
AA-Wood-4        0.33
ChipMoisture-4   0.33
SteamFlow-4      0.33
Lower-HeatT-3    0.33
Upper-HeatT-3    0.33
ChipMass-4       0.33
WeakLiquorF      0.33
BlackFlow-2      0.33
WeakWashF        0.33
SteamHeatF-3     0.33
T-Top-Chips-4    0.33
SulphidityL-4    46.84
dtype: float64
```

As expected, the datapoint 'AAWhiteSt-4' even has 46% of data missing! It seems we only have 300 datapoints and presumably these missing values occur in different records our dataset will be decimated if we just drop all rows with missing values.

```
kamyr_df.shape
```

```
(301, 23)
```

```
kamyr_df.dropna().shape
```

```
(131, 23)
```

As we drop all rows with missing values, we are left with only 131 records. Whilst this might be good enough for some purposes, there are more viable options.

Perhaps we can first remove the column with the most missing values and then drop all remaining

```
kamyr_df.drop(columns=['AAWhiteSt-4 ', 'SulphidityL-4 ']).dropna().shape
```



```
(263, 21)
```

Significantly better, although we lost the information of 2 sensors we now have a complete dataset with 263 records. For purposes where those 2 sensors are irrelevant this is a viable option, keep in mind that this dataset is still 100% truthful, as we have not imputed any values.

Another option, where we retain all our records would be using the timely nature of our dataset, each record is a measurement with an interval of 1 hour. I have no knowledge of this dataset but one might make the assumption that the interval of 1 hour is taken as the state of the machine does not alter much in 1 hour. Therefore we could do what is called a forward fill, where we fill in the missing values with the same value of the sensor for the previous measurement.

This would solve nearly all nan values as there might be a problem where the first value is missing. This is shown below.

```
kamyr_df.fillna(method='ffill')['SulphidityL-4 ']
```

```
0      NaN
1    29.11
2    29.11
3    29.02
4    29.01
...
296   30.43
297   30.29
298   30.47
299   30.47
300   30.46
Name: SulphidityL-4 , Length: 301, dtype: float64
```

Although our dataset is not fully the truth, we can see that little to no changes occur in the sensor and using a forward fill is arguably the most suitable option.

3.2 Travel times

Another dataset from the same source contains a collection of recorded travel times and specific information about the travel itself as e.g.: the day of the week, where they were going, ...

```
travel_df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-
practical-approach/main/src/c2_data_preparation/data/travel-times.csv')
travel_df
```

	Date	StartTime	DayOfWeek	GoingTo	Distance	MaxSpeed	AvgSpeed	\
0	1/6/2012	16:37	Friday	Home	51.29	127.4	78.3	
1	1/6/2012	08:20	Friday	GSK	51.63	130.3	81.8	
2	1/4/2012	16:17	Wednesday	Home	51.27	127.4	82.0	
3	1/4/2012	07:53	Wednesday	GSK	49.17	132.3	74.2	
4	1/3/2012	18:57	Tuesday	Home	51.15	136.2	83.4	
..	
200	7/18/2011	08:09	Monday	GSK	54.52	125.6	49.9	
201	7/14/2011	08:03	Thursday	GSK	50.90	123.7	76.2	
202	7/13/2011	17:08	Wednesday	Home	51.96	132.6	57.5	
203	7/12/2011	17:51	Tuesday	Home	53.28	125.8	61.6	
204	7/11/2011	16:56	Monday	Home	51.73	125.0	62.8	
AvgMovingSpeed FuelEconomy TotalTime MovingTime Take407All Comments								

(continues on next page)

(continued from previous page)

0	84.8	NaN	39.3	36.3	No	NaN
1	88.9	NaN	37.9	34.9	No	NaN
2	85.8	NaN	37.5	35.9	No	NaN
3	82.9	NaN	39.8	35.6	No	NaN
4	88.1	NaN	36.8	34.8	No	NaN
..
200	82.4	7.89	65.5	39.7	No	NaN
201	95.1	7.89	40.1	32.1	Yes	NaN
202	76.7	NaN	54.2	40.6	Yes	NaN
203	87.6	NaN	51.9	36.5	Yes	NaN
204	92.5	NaN	49.5	33.6	Yes	NaN
[205 rows x 13 columns]						

we have a total of 205 records and we can already see that the FuelEconomy column seems pretty bad, let's quantify that.

```
travel_df.isna().sum().divide(len(travel_df)).round(4)*100
```

```
Date          0.00
StartTime      0.00
DayOfWeek      0.00
GoingTo        0.00
Distance       0.00
MaxSpeed       0.00
AvgSpeed       0.00
AvgMovingSpeed 0.00
FuelEconomy    8.29
TotalTime      0.00
MovingTime     0.00
Take407All     0.00
Comments      88.29
dtype: float64
```

In the end, it doesn't seem that bad, but there are comments and nearly none of them are filled in. Which in perspective is understandable. Let's see what the comments look like

```
travel_df[~travel_df.Comments.isna()].Comments
```

```
15          Put snow tires on
39          Heavy rain
49      Huge traffic backup
50      Pumped tires up: check fuel economy improved?
52          Backed up at Bronte
54          Backed up at Bronte
60          Rainy
78          Rain, rain, rain
91          Rain, rain, rain
92      Accident: backup from Hamilton to 407 ramp
110         Raining
132         Back to school traffic?
133      Took 407 all the way (to McMaster)
150         Heavy volume on Derry
156         Start early to run a batch
158      Accident at 403/highway 6; detour along Dundas
165         Detour taken
166         Must be Friday
```

(continues on next page)

(continued from previous page)

```

172             Medium amount of rain
174             New tires
182             Turn around on Derry
184             Empty roads
187             Police slowdown on 403
189             Accident blocked 407 exit
Name: Comments, dtype: object

```

As you would expect, these comments are text based. Now imagine we would like to run some Natural Language Processing (NLP) on these, it would be a pain to perform string operations on it when it is riddled with missing values.

Here a simple example where we select all records containing the word 'rain', with no avail.

```
travel_df[travel_df.Comments.str.lower().str.contains('rain')]
```

```

-----
ValueError                                Traceback (most recent call last)
/tmp/ipykernel_25543/1298831137.py in <module>
----> 1 travel_df[travel_df.Comments.str.lower().str.contains('rain')]

~/git/data-science-practical-approach/venv/lib/python3.8/site-packages/pandas/core/
frame.py in __getitem__(self, key)
   3446
   3447         # Do we have a (boolean) 1d indexer?
-> 3448         if com.is_bool_indexer(key):
   3449             return self._getitem_bool_array(key)
   3450

~/git/data-science-practical-approach/venv/lib/python3.8/site-packages/pandas/core/
common.py in is_bool_indexer(key)
   137         # Don't raise on e.g. ["A", "B", np.nan], see
   138         # test_loc_getitem_list_of_labels_categoricalindex_with_
-> na
-> 139         raise ValueError(na_msg)
   140         return False
   141         return True

ValueError: Cannot mask with non-boolean array containing NA / NaN values

```

The last line of the python error traceback gives us the reason it failed, because there were NaN values present.

Luckily the string variable has more or less it's on 'null' value, being an empty string, this way these operations are still possible, most of the comments will just contain nothing.

```
travel_df.Comments = travel_df.Comments.fillna('')
```

```
travel_df[travel_df.Comments.str.lower().str.contains('rain')]
```

	Date	StartTime	DayOfWeek	GoingTo	Distance	MaxSpeed	AvgSpeed	\
39	11/29/2011	07:23	Tuesday	GSK	51.74	112.2	55.3	
60	11/9/2011	16:15	Wednesday	Home	51.28	121.4	65.9	
78	10/25/2011	17:24	Tuesday	Home	52.87	123.5	65.1	
91	10/12/2011	17:47	Wednesday	Home	51.40	114.4	59.7	
110	9/27/2011	07:36	Tuesday	GSK	50.65	128.1	86.3	
172	8/9/2011	08:15	Tuesday	GSK	49.08	134.8	60.5	

(continues on next page)

(continued from previous page)

	AvgMovingSpeed	FuelEconomy	TotalTime	MovingTime	Take407All	\
39	61.0	NaN	56.2	50.9	No	
60	71.8	9.35	46.7	42.1	No	
78	72.4	8.97	48.7	43.8	No	
91	65.8	8.75	51.7	46.9	No	
110	88.6	8.31	35.2	34.3	Yes	
172	67.2	8.54	48.7	43.8	No	
	Comments					
39	Heavy rain					
60	Rainy					
78	Rain, rain, rain					
91	Rain, rain, rain					
110	Raining					
172	Medium amount of rain					

Fixed! now we can use the comments for analysis.

We still have to fix the FuelEconomy, let us take a look at the non NaN values

```
travel_df[~travel_df.FuelEconomy.isna()]
```

	Date	StartTime	DayOfWeek	GoingTo	Distance	MaxSpeed	AvgSpeed	\
6	1/2/2012	17:31	Monday	Home	51.37	123.2	82.9	
7	1/2/2012	07:34	Monday	GSK	49.01	128.3	77.5	
8	12/23/2011	08:01	Friday	GSK	52.91	130.3	80.9	
9	12/22/2011	17:19	Thursday	Home	51.17	122.3	70.6	
10	12/22/2011	08:16	Thursday	GSK	49.15	129.4	74.0	
..	
197	7/20/2011	08:24	Wednesday	GSK	48.50	125.8	75.7	
198	7/19/2011	17:17	Tuesday	Home	51.16	126.7	92.2	
199	7/19/2011	08:11	Tuesday	GSK	50.96	124.3	82.3	
200	7/18/2011	08:09	Monday	GSK	54.52	125.6	49.9	
201	7/14/2011	08:03	Thursday	GSK	50.90	123.7	76.2	
	AvgMovingSpeed	FuelEconomy	TotalTime	MovingTime	Take407All	Comments		
6	87.3	-	37.2	35.3	No			
7	85.9	-	37.9	34.3	No			
8	88.3	8.89	39.3	36.0	No			
9	78.1	8.89	43.5	39.3	No			
10	81.4	8.89	39.8	36.2	No			
..		
197	87.3	7.89	38.5	33.3	Yes			
198	102.6	7.89	33.3	29.9	Yes			
199	96.4	7.89	37.2	31.7	Yes			
200	82.4	7.89	65.5	39.7	No			
201	95.1	7.89	40.1	32.1	Yes			

[188 rows x 13 columns]

It seems that aside NaN values there are also other intruders, a quick check on the data type (Dtype) reveals it is not recognised as a number!

```
travel_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   205 non-null   object
1   StartTime              205 non-null   object
2   DayOfWeek              205 non-null   object
3   GoingTo                205 non-null   object
4   Distance               205 non-null   float64
5   MaxSpeed              205 non-null   float64
6   AvgSpeed              205 non-null   float64
7   AvgMovingSpeed         205 non-null   float64
8   FuelEconomy            188 non-null   object
9   TotalTime             205 non-null   float64
10  MovingTime            205 non-null   float64
11  Take407All            205 non-null   object
12  Comments              205 non-null   object
dtypes: float64(6), object(7)
memory usage: 20.9+ KB
```

The column is noted as an object or string type, meaning that these numbers are given as '9.24' instead of 9.24 and numerical operations are not possible. We can cast them to numeric but have to warn pandas to coerce errors, meaning errors will be converted to NaN values. Later we'll handle the NaN's.

```
travel_df.FuelEconomy = pd.to_numeric(travel_df.FuelEconomy, errors='coerce')
travel_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   205 non-null   object
1   StartTime              205 non-null   object
2   DayOfWeek              205 non-null   object
3   GoingTo                205 non-null   object
4   Distance               205 non-null   float64
5   MaxSpeed              205 non-null   float64
6   AvgSpeed              205 non-null   float64
7   AvgMovingSpeed         205 non-null   float64
8   FuelEconomy            186 non-null   float64
9   TotalTime             205 non-null   float64
10  MovingTime            205 non-null   float64
11  Take407All            205 non-null   object
12  Comments              205 non-null   object
dtypes: float64(7), object(6)
memory usage: 20.9+ KB
```

Wonderful, now the column is numerical and we can see 2 more missing values have popped up! We could easily drop these 19 records and have a complete dataset.

```
travel_df.dropna()
```

	Date	StartTime	DayOfWeek	GoingTo	Distance	MaxSpeed	AvgSpeed	\
8	12/23/2011	08:01	Friday	GSK	52.91	130.3	80.9	

(continues on next page)

(continued from previous page)

9	12/22/2011	17:19	Thursday	Home	51.17	122.3	70.6
10	12/22/2011	08:16	Thursday	GSK	49.15	129.4	74.0
11	12/21/2011	07:45	Wednesday	GSK	51.77	124.8	71.7
12	12/20/2011	16:05	Tuesday	Home	51.45	130.1	75.2
..
197	7/20/2011	08:24	Wednesday	GSK	48.50	125.8	75.7
198	7/19/2011	17:17	Tuesday	Home	51.16	126.7	92.2
199	7/19/2011	08:11	Tuesday	GSK	50.96	124.3	82.3
200	7/18/2011	08:09	Monday	GSK	54.52	125.6	49.9
201	7/14/2011	08:03	Thursday	GSK	50.90	123.7	76.2
	AvgMovingSpeed	FuelEconomy	TotalTime	MovingTime	Take407All	Comments	
8	88.3	8.89	39.3	36.0	No		
9	78.1	8.89	43.5	39.3	No		
10	81.4	8.89	39.8	36.2	No		
11	78.9	8.89	43.3	39.4	No		
12	82.7	8.89	41.1	37.3	No		
..
197	87.3	7.89	38.5	33.3	Yes		
198	102.6	7.89	33.3	29.9	Yes		
199	96.4	7.89	37.2	31.7	Yes		
200	82.4	7.89	65.5	39.7	No		
201	95.1	7.89	40.1	32.1	Yes		
[186 rows x 13 columns]							

However im leaving them as an exercise for you to apply a technique we will see in the next part

3.3 Material properties

Another dataset from the same source contains the material properties from 30 samples, this time there is not timestamp as the samples are not related in time with each other.

```
material_df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-
practical-approach/main/src/c2_data_preparation/data/raw-material-properties.csv')
material_df
```

	Sample	size1	size2	size3	density1	density2	density3
0	X12558	0.696	2.69	6.38	41.8	17.18	3.90
1	X14728	0.636	2.30	5.14	38.1	12.73	3.89
2	X15468	0.841	2.85	5.20	37.6	13.58	3.98
3	X21364	0.609	2.13	4.62	34.2	11.12	4.02
4	X23671	0.684	2.16	4.87	36.4	12.24	3.92
5	X24055	0.762	2.81	6.36	38.1	13.28	3.89
6	X24905	0.552	2.34	5.03	41.3	16.71	3.86
7	X25917	0.501	2.17	5.09	NaN	NaN	NaN
8	X27871	0.619	2.11	5.13	NaN	NaN	NaN
9	X28690	0.610	2.10	4.18	35.0	12.15	3.86
10	X31385	0.532	2.09	4.93	NaN	NaN	NaN
11	X31813	0.738	2.29	5.47	NaN	NaN	NaN
12	X32807	0.779	2.62	5.59	NaN	NaN	NaN
13	X33943	0.537	2.23	5.41	35.2	11.34	3.99
14	X35035	0.702	2.05	5.10	34.2	10.54	4.02
15	X39223	0.768	2.51	5.09	34.9	12.55	3.90

(continues on next page)

(continued from previous page)

16	X40503	0.714	2.56	6.03	35.6	12.20	4.02
17	X41400	0.621	2.42	5.10	38.7	14.27	3.98
18	X42988	0.726	2.11	4.69	37.1	13.14	3.98
19	X44749	0.698	2.36	5.40	36.6	12.16	4.01
20	X45295	NaN	NaN	NaN	38.1	13.34	3.89
21	X46965	0.759	2.47	4.83	38.7	14.83	3.89
22	X49666	0.535	2.13	5.23	NaN	NaN	NaN
23	X50678	0.716	2.29	5.45	37.3	13.70	3.92
24	X52894	0.635	2.08	4.94	NaN	NaN	NaN
25	X53925	0.598	2.12	4.69	37.9	13.45	3.78
26	X54254	0.700	2.47	5.22	38.8	14.72	3.92
27	X54272	0.957	2.96	7.37	36.2	13.38	4.20
28	X54394	0.759	2.66	5.36	35.2	12.19	3.98
29	X55408	0.661	2.10	4.27	NaN	NaN	NaN
30	X56952	0.646	2.38	4.51	40.1	15.68	3.86
31	X57095	0.662	2.34	4.71	35.0	12.37	3.90
32	X57128	0.749	2.43	5.16	37.3	13.04	3.92
33	X61870	0.598	2.21	4.90	NaN	NaN	NaN
34	X61888	0.619	2.59	5.81	NaN	NaN	NaN
35	X72736	0.693	2.05	5.02	39.6	15.55	3.94

let us quantify the amount of missing data

```
material_df.isna().sum().divide(len(material_df)).round(4)*100
```

```
Sample      0.00
size1       2.78
size2       2.78
size3       2.78
density1    27.78
density2    27.78
density3    27.78
dtype: float64
```

Unfortunately that is a lot of missing data, covered in all records, dropping here seems almost impossible if we want to keep a healthy amount of records.

Here it would be wise to go for a more elaborate method of imputation, I opted for the K-nearest neighbours method, which looks at the K most similar records in the dataset to make an educated guess on what the missing value could be, this because we can assume that records with similar data are also similar over all the properties (columns).

Im using the sklearn library for this, which has more imputation techniques such as MICE. More info can be found [here](#)

```
from sklearn.impute import KNNImputer
```

im creating an imputer object and specify that i want to use the 5 most similar records and weigh them by distance from the to imputed record, meaning closer neighbours are more important.

```
imputer = KNNImputer(n_neighbors=5, weights="distance")
```

As the imputer only takes numerical values I had to do some pandas magic and drop the first column, which I then added again. The result is a fully filled dataset, you can recognise the new values as they are not rounded.

```
pd.DataFrame(
    imputer.fit_transform(material_df.drop(columns=['Sample'])),
    columns=material_df.columns.drop('Sample')
)
```

	size1	size2	size3	density1	density2	density3
0	0.696000	2.690000	6.380000	41.800000	17.180000	3.900000
1	0.636000	2.300000	5.140000	38.100000	12.730000	3.890000
2	0.841000	2.850000	5.200000	37.600000	13.580000	3.980000
3	0.609000	2.130000	4.620000	34.200000	11.120000	4.020000
4	0.684000	2.160000	4.870000	36.400000	12.240000	3.920000
5	0.762000	2.810000	6.360000	38.100000	13.280000	3.890000
6	0.552000	2.340000	5.030000	41.300000	16.710000	3.860000
7	0.501000	2.170000	5.090000	38.495282	14.029399	3.931180
8	0.619000	2.110000	5.130000	37.405275	13.157346	3.943667
9	0.610000	2.100000	4.180000	35.000000	12.150000	3.860000
10	0.532000	2.090000	4.930000	37.811132	13.646072	3.908364
11	0.738000	2.290000	5.470000	37.088833	13.255412	3.941654
12	0.779000	2.620000	5.590000	36.540567	12.889902	3.970973
13	0.537000	2.230000	5.410000	35.200000	11.340000	3.990000
14	0.702000	2.050000	5.100000	34.200000	10.540000	4.020000
15	0.768000	2.510000	5.090000	34.900000	12.550000	3.900000
16	0.714000	2.560000	6.030000	35.600000	12.200000	4.020000
17	0.621000	2.420000	5.100000	38.700000	14.270000	3.980000
18	0.726000	2.110000	4.690000	37.100000	13.140000	3.980000
19	0.698000	2.360000	5.400000	36.600000	12.160000	4.010000
20	0.733097	2.653959	5.881504	38.100000	13.340000	3.890000
21	0.759000	2.470000	4.830000	38.700000	14.830000	3.890000
22	0.535000	2.130000	5.230000	37.391815	13.089536	3.944335
23	0.716000	2.290000	5.450000	37.300000	13.700000	3.920000
24	0.635000	2.080000	4.940000	37.254724	13.206262	3.933904
25	0.598000	2.120000	4.690000	37.900000	13.450000	3.780000
26	0.700000	2.470000	5.220000	38.800000	14.720000	3.920000
27	0.957000	2.960000	7.370000	36.200000	13.380000	4.200000
28	0.759000	2.660000	5.360000	35.200000	12.190000	3.980000
29	0.661000	2.100000	4.270000	36.172345	12.755632	3.887375
30	0.646000	2.380000	4.510000	40.100000	15.680000	3.860000
31	0.662000	2.340000	4.710000	35.000000	12.370000	3.900000
32	0.749000	2.430000	5.160000	37.300000	13.040000	3.920000
33	0.598000	2.210000	4.900000	37.865882	13.826029	3.887021
34	0.619000	2.590000	5.810000	35.932339	12.318210	3.989911
35	0.693000	2.050000	5.020000	39.600000	15.550000	3.940000

This concludes the part of missing values, perhaps you can try yourself and impute the missing values for the FuelEconomy using the SimpleImputer or even the IterativeImputer.

CONCATENATION AND DEDUPLICATION

In this notebook we are going to investigate the concepts of stitching data files (concatenation) and verifying the integrity of our data concerning duplicates

4.1 Concatenation

When dealing with large amounts of data, fractioning is often the only solution. Not only does this tidy up your data space, but it also benefits computation. Aside from that, appending new data to your data lake is independent of the historical data. However if you want to perform historical analysis this means you will need to perform additional operations.

In this notebook we have a setup of a very small data lake containing daily minimal temperatures. If you would look closely in the url you would see the following structure.

data/temperature/australia/melbourne/1981.csv

This is a straight-forward but perfect example on how fragmentation works, in our data lake we have: temperatures data fractioned by country, city and year. As we are working with daily temperatures further fractioning would not be interesting, but you could fraction e.g. per month.

In the cells below, we read our both 1981 and 1982 data and concatenate them using python.

```
import pandas as pd
```

```
melbourne_1981_df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-practical-approach/main/src/c2_data_preparation/data/temperatures/australia/melbourne/1981.csv')
```

```
melbourne_1982_df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-practical-approach/main/src/c2_data_preparation/data/temperatures/australia/melbourne/1982.csv')
```

```
df = pd.concat([
    melbourne_1981_df,
    melbourne_1982_df,
])
```

```
df
```

```
      Date  Temp
0  1981-01-01  20.7
1  1981-01-02  17.9
2  1981-01-03  18.8
3  1981-01-04  14.6
4  1981-01-05  15.8
..      ...    ...
360 1982-12-27  15.3
361 1982-12-28  16.3
362 1982-12-29  15.8
363 1982-12-30  17.7
364 1982-12-31  16.3
```

```
[730 rows x 2 columns]
```

And there you have it! we now have a dataframe containing both data from 1981 as 1982. Can you figure out what I calculated in the next cell? Do you think there might be a more ‘clean’ solution?

```
df[df.Date.str[5:7]=='01'].Temp.mean()
```

```
17.140322580645158
```

As an exercise I would ask you now to create a small python script that given a begin and end year (between 1981 and 1990) can automatically concatenate all the necessary data

```
for i in range(1982,1987):
    print(i)
```

```
1982
1983
1984
1985
1986
```

4.2 Deduplication

Another important aspect of data cleaning is the removal of duplicates. Here we fragment of a dataset from activity on a popular games platform. We can see which user has either bought or played specific games and how often. Unfortunately for some reason, entries might have duplicates which we have to deal with as otherwise users might have e.g. bought a game twice.

```
df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-practical-
->approach/main/src/c2_data_preparation/data/steam.csv')
df
```

```
   user_id      game      action  freq
0  11373749  Sid Meier's Civilization IV  purchase    1.0
1  11373749  Sid Meier's Civilization IV    play    0.1
2  11373749  Sid Meier's Civilization IV  purchase    1.0
3  11373749  Sid Meier's Civilization IV Beyond the Sword  purchase    1.0
4  11373749  Sid Meier's Civilization IV Beyond the Sword  purchase    1.0
...      ...      ...      ...    ...
1834 112845094      Arma 2  purchase    1.0
```

(continues on next page)

(continued from previous page)

1835	112845094	Grand Theft Auto San Andreas	purchase	1.0
1836	112845094	Grand Theft Auto Vice City	purchase	1.0
1837	112845094	Grand Theft Auto Vice City	purchase	1.0
1838	112845094	Grand Theft Auto III	purchase	1.0

[1839 rows x 4 columns]

We have a dataframe with 1839 interactions, you can see that the freq either notes the amount they bought (which always 1 as there is not use in buying it more) or the amount in hours they played.

Let us straightforward ask pandas to remove all rows that have an exact duplicate

```
df.drop_duplicates()
```

	user_id	game	action	freq
0	11373749	Sid Meier's Civilization IV	purchase	1.0
1	11373749	Sid Meier's Civilization IV	play	0.1
3	11373749	Sid Meier's Civilization IV Beyond the Sword	purchase	1.0
5	11373749	Sid Meier's Civilization IV Warlords	purchase	1.0
7	56038151	Tom Clancy's H.A.W.X. 2	purchase	1.0
...
1831	112845094	Grand Theft Auto San Andreas	purchase	1.0
1832	112845094	Grand Theft Auto San Andreas	play	0.2
1833	112845094	Grand Theft Auto III	purchase	1.0
1834	112845094	Arma 2	purchase	1.0
1836	112845094	Grand Theft Auto Vice City	purchase	1.0

[1132 rows x 4 columns]

Alright! this seemed to have dropped 707 rows from our dataset, but we would like to know more about those. Let's ask which rows the algorithm has dropped:

```
df[df.duplicated()]
```

	user_id	game	action	freq
2	11373749	Sid Meier's Civilization IV	purchase	1.0
4	11373749	Sid Meier's Civilization IV Beyond the Sword	purchase	1.0
6	11373749	Sid Meier's Civilization IV Warlords	purchase	1.0
10	56038151	Grand Theft Auto San Andreas	purchase	1.0
12	56038151	Grand Theft Auto Vice City	purchase	1.0
...
1827	39146470	Sid Meier's Civilization IV Warlords	purchase	1.0
1830	48666962	Crysis 2	purchase	1.0
1835	112845094	Grand Theft Auto San Andreas	purchase	1.0
1837	112845094	Grand Theft Auto Vice City	purchase	1.0
1838	112845094	Grand Theft Auto III	purchase	1.0

[707 rows x 4 columns]

Here we can see the duplicates, no particular pattern seems to be present, we could just for curiosity count the games that are duplicated

```
df[df.duplicated()].game.value_counts()
```

Grand Theft Auto San Andreas	172
Grand Theft Auto Vice City	103

(continues on next page)

(continued from previous page)

Sid Meier's Civilization IV	98
Grand Theft Auto III	90
Sid Meier's Civilization IV Beyond the Sword	80
Sid Meier's Civilization IV Warlords	79
Sid Meier's Civilization IV Colonization	75
Crysis 2	7
Arma 2	1
Tom Clancy's H.A.W.X. 2	1
TERA	1
Name: game, dtype: int64	

It seems there are some games which are very prone to being duplicated, at this point we could go and ask the IT department why these games are acting weird.

Another thing im interested about is the perspective of a single gamer, here we took a single user_id and printed all his games

```
df[df.user_id == 11373749]
```

	user_id	game	action	freq
0	11373749	Sid Meier's Civilization IV	purchase	1.0
1	11373749	Sid Meier's Civilization IV	play	0.1
2	11373749	Sid Meier's Civilization IV	purchase	1.0
3	11373749	Sid Meier's Civilization IV Beyond the Sword	purchase	1.0
4	11373749	Sid Meier's Civilization IV Beyond the Sword	purchase	1.0
5	11373749	Sid Meier's Civilization IV Warlords	purchase	1.0
6	11373749	Sid Meier's Civilization IV Warlords	purchase	1.0

Ah, you can see all of his three games are somehow duplicated in purchase, also it seems he only played one of them for only 0.1 hours. Looks like he fell to the bait of a tempting summer sale but didn't realise he had no time to actually play it.

Another thing I would like to mention here is that this dataset would make a fine recommender system as it contains user ids and hours played. Add game metadata (description) and reviews to the mix and your data preparation is done!

We can remove all duplicates now by overwriting our dataframe

```
df = df.drop_duplicates()
```

One thing still bothers me, as hours played can change over time it might be that different snapshots have produced different values, therefore more duplicates might be present with different hours_played.

Time to investigate this by using a subset of columns in the drop_duplicates algorithm

```
df.drop_duplicates(subset=['user_id', 'game', 'action'])
```

	user_id	game	action	freq
0	11373749	Sid Meier's Civilization IV	purchase	1.0
1	11373749	Sid Meier's Civilization IV	play	0.1
3	11373749	Sid Meier's Civilization IV Beyond the Sword	purchase	1.0
5	11373749	Sid Meier's Civilization IV Warlords	purchase	1.0
7	56038151	Tom Clancy's H.A.W.X. 2	purchase	1.0
...
1831	112845094	Grand Theft Auto San Andreas	purchase	1.0
1832	112845094	Grand Theft Auto San Andreas	play	0.2
1833	112845094	Grand Theft Auto III	purchase	1.0
1834	112845094	Arma 2	purchase	1.0

(continues on next page)

(continued from previous page)

```
1836 112845094          Grand Theft Auto Vice City  purchase    1.0

[1120 rows x 4 columns]
```

Seems we have shaved off another 12 records, so our intuition was right, again lets see which the duplicates are:

```
df[df.duplicated(subset=['user_id', 'game', 'action'])]
```

	user_id	game	action	freq
118	118664413	Grand Theft Auto San Andreas	play	0.2
458	50769696	Grand Theft Auto San Andreas	play	3.1
521	71411882	Grand Theft Auto III	play	0.2
607	33865373	Sid Meier's Civilization IV	play	2.0
898	71510748	Grand Theft Auto San Andreas	play	0.2
908	28472068	Grand Theft Auto Vice City	play	0.4
910	28472068	Grand Theft Auto San Andreas	play	0.2
912	28472068	Grand Theft Auto III	play	0.1
1506	59925638	Tom Clancy's H.A.W.X. 2	play	0.3
1553	148362155	Grand Theft Auto San Andreas	play	12.5
1709	176261926	Sid Meier's Civilization IV Beyond the Sword	play	0.4
1711	176261926	Sid Meier's Civilization IV	play	0.2

As expected the duplicates are all in the 'play' action, to complete our view we extract the data of a single user

```
df[df.user_id==118664413]
```

	user_id	game	action	freq
115	118664413	Grand Theft Auto San Andreas	purchase	1.0
116	118664413	Grand Theft Auto San Andreas	play	1.9
118	118664413	Grand Theft Auto San Andreas	play	0.2

It looks like we have a problem now, we know these are duplicates and should be removed, but which one? Personally I would argue here that we keep the highest value, as it is impossible to 'unplay' hours on the game. I will leave this as an exercise for you, but the solution is pretty tricky so i'll give a hint:

The algorithm always keeps the first record in case of duplicates, so you could sort the rows making sure the higher value is always encountered first, good luck!

OUTLIERS AND VALIDITY

```
import pandas as pd
```

```
wafer_df = pd.read_csv('https://openmv.net/file/silicon-wafer-thickness.csv')  
wafer_df.head()
```

	G1	G2	G3	G4	G5	G6	G7	G8	G9
0	0.175	0.188	-0.159	0.095	0.374	-0.238	-0.800	0.158	-0.211
1	0.102	0.075	0.141	0.180	0.138	-0.057	-0.075	0.072	0.072
2	0.607	0.711	0.879	0.765	0.592	0.187	0.431	0.345	0.187
3	0.774	0.823	0.619	0.370	0.725	0.439	-0.025	-0.259	0.496
4	0.504	0.644	0.845	0.681	0.502	0.151	0.404	0.296	0.260

```
iqr = wafer_df.quantile(0.75)-wafer_df.quantile(0.25)
```

```
range_df = (wafer_df-wafer_df.quantile(0.5))/iqr
```

```
range_df[(range_df>2).any(axis='columns')]
```

	G1	G2	G3	G4	G5	G6 \
8	2.232430	2.009016	1.956542	1.589328	1.843890	1.544669
38	12.891135	12.827049	12.832178	13.913292	11.429506	9.500865
39	3.691318	3.981148	3.774387	4.081944	3.248059	3.729107
61	2.010106	2.153279	1.987980	1.863745	1.858602	1.274928
110	3.678457	2.841803	3.204808	3.180562	2.669391	0.518732
112	2.361047	2.086066	2.363384	2.107670	1.925623	1.238040
117	1.475425	1.043443	2.154415	2.582182	0.653862	1.823631
120	1.791456	1.484426	2.583449	1.440686	2.085819	0.990202
121	1.791456	1.484426	2.583449	1.440686	2.085819	0.990202
152	2.610932	2.102459	2.387425	2.549786	2.169187	1.730259
154	-0.529169	-0.538525	-0.404993	-0.331586	-0.552513	4.565994

	G7	G8	G9
8	1.233344	0.419604	1.582851
38	10.305875	9.927200	9.055620
39	3.304890	3.846374	3.149479
61	1.237283	0.825451	0.955968
110	0.700361	0.176555	0.727694
112	1.766328	0.890800	1.377752
117	1.581227	0.857552	1.188876
120	1.782081	1.034107	1.822711
121	1.782081	1.034107	1.822711

(continues on next page)

(continued from previous page)

```
152  2.241549  1.713958  1.592121
154  -0.051854 -0.382918 -0.536501
```

```
range_df[(range_df<=-2).any(axis='columns')]
```

```

      G1      G2      G3      G4      G5      G6      G7  \
54 -1.550758 -1.525410 -1.843736 -2.082897 -1.659174 -1.203458 -1.184772
56 -1.732660 -1.510656 -2.121128 -2.122916 -1.781774 -1.521614 -1.909419
59 -1.971520 -1.310656 -2.328248 -1.175798 -2.067838 -0.915274 -1.783394
64 -1.234727 -1.361475 -0.736015 -1.055741 -2.224765 -0.839193 -0.679357
65 -2.226918 -1.194262 -2.117429 -2.161029 -2.043318 -0.190202 -1.004923
102 -2.484153 -2.330328 -1.568192 -2.808957 -1.945239 -1.340634 -0.846078

      G8      G9
54 -1.650903 -1.245655
56 -1.782746 -1.159907
59 -1.304672 -1.514484
64 -0.865578 -0.663963
65 -0.270565 -0.794902
102 -1.691029 -0.887601
```

```
from sklearn.ensemble import IsolationForest
```

```
clf = IsolationForest(random_state=0).fit(wafer_df)
wafer_df[clf.predict(wafer_df)==-1]
```

```

      G1      G2      G3      G4      G5      G6      G7      G8      G9
8    1.396  1.461  1.342  1.122  1.394  1.408  0.924  0.638  1.375
20  -0.558 -0.705 -0.526 -0.412 -0.753 -0.998 -0.270  0.598 -1.416
38   7.197  8.060  7.223  7.589  7.258  8.310  7.835  8.931  7.824
39   2.190  2.664  2.325  2.430  2.253  3.303  2.502  3.627  2.727
54  -0.663 -0.695 -0.713 -0.805 -0.749 -0.976 -0.918 -1.168 -1.066
56  -0.762 -0.686 -0.863 -0.826 -0.824 -1.252 -1.470 -1.283 -0.992
59  -0.892 -0.564 -0.975 -0.329 -0.999 -0.726 -1.374 -0.866 -1.298
61   1.275  1.549  1.359  1.266  1.403  1.174  0.927  0.992  0.834
65  -1.031 -0.493 -0.861 -0.846 -0.984 -0.097 -0.781  0.036 -0.677
102 -1.171 -1.186 -0.564 -1.186 -0.924 -1.095 -0.660 -1.203 -0.757
106 -0.659 -0.451 -0.692 -0.708 -0.595 -0.726 -1.031 -0.877 -1.080
110  2.183  1.969  2.017  1.957  1.899  0.518  0.518  0.426  0.637
112  1.466  1.508  1.562  1.394  1.444  1.142  1.330  1.049  1.198
117  0.984  0.872  1.449  1.643  0.666  1.650  1.189  1.020  1.035
120  1.156  1.141  1.681  1.044  1.542  0.927  1.342  1.174  1.582
121  1.156  1.141  1.681  1.044  1.542  0.927  1.342  1.174  1.582
152  1.602  1.518  1.575  1.626  1.593  1.569  1.692  1.767  1.383
```


SOME PRACTICE

Now that you have learned techniques in data preparation, why don't you put them to use in this wonderfully horrifying dataset. Good luck!

```
import os
import json

import pandas as pd
```

```
kaggle_dir = os.path.expanduser("~/kaggle")
if not os.path.exists(kaggle_dir):
    os.mkdir(kaggle_dir)

with open(f'{kaggle_dir}/kaggle.json', 'w') as f:
    json.dump(
        {
            "username": "lorenzof",
            "key": "7a44a9e99b27e796177d793a3d85b8cf"
        }, f)
```

```
import kaggle
kaggle.api.dataset_download_files(dataset='PromptCloudHQ/us-jobs-on-monstercom', path=
↳ './data', unzip=True)
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
/tmp/ipykernel_25600/39646943.py in <module>
----> 1 import kaggle
      2 kaggle.api.dataset_download_files(dataset='PromptCloudHQ/us-jobs-on-monstercom
↳ ', path='./data', unzip=True)

ModuleNotFoundError: No module named 'kaggle'
```

```
df = pd.read_csv('./data/monster_com-job_sample.csv')
```

```
df.head()
```

	country	country_code	date_added	has_expired	\
0	United States of America	US	NaN	No	
1	United States of America	US	NaN	No	
2	United States of America	US	NaN	No	

(continues on next page)

(continued from previous page)

3	United States of America	US	NaN	No
4	United States of America	US	NaN	No
	job_board		job_description	\
0	jobs.monster.com	TeamSoft is seeing an IT Support Specialist to...		
1	jobs.monster.com	The Wisconsin State Journal is seeking a flexi...		
2	jobs.monster.com	Report this job About the Job DePuy Synthes Co...		
3	jobs.monster.com	Why Join Altec? If you're considering a career...		
4	jobs.monster.com	Position ID# 76162 # Positions 1 State CT C...		
		job_title	job_type	\
0	IT Support Technician Job in Madison	Full Time Employee		
1	Business Reporter/Editor Job in Madison	Full Time		
2	Johnson & Johnson Family of Companies Job Appl...	Full Time, Employee		
3	Engineer - Quality Job in Dixon	Full Time		
4	Shift Supervisor - Part-Time Job in Camphill	Full Time Employee		
		location		\
0		Madison, WI 53702		
1		Madison, WI 53708		
2	DePuy Synthes Companies is a member of Johnson...			
3		Dixon, CA		
4		Camphill, PA		
	organization			\
0		NaN		
1	Printing and Publishing			
2	Personal and Household Services			
3	Altec Industries			
4	Retail			
		page_url	salary	\
0	http://jobview.monster.com/it-support-technici...	NaN		
1	http://jobview.monster.com/business-reporter-e...	NaN		
2	http://jobview.monster.com/senior-training-lea...	NaN		
3	http://jobview.monster.com/engineer-quality-jo...	NaN		
4	http://jobview.monster.com/shift-supervisor-pa...	NaN		
	sector		uniq_id	
0	IT/Software Development	11d599f229a80023d2f40e7c52cd941e		
1		NaN	e4cbb126dabf22159aff90223243ff2a	
2		NaN	839106b353877fa3d896ffb9c1fe01c0	
3	Experienced (Non-Manager)	58435fcab804439efdcaa7ecca0fd783		
4	Project/Program Management	64d0272dc8496abfd9523a8df63c184c		

Need some inspiration? perhaps [this](#) might help!

Part III

3. Data Preprocessing

DATA PREPROCESSING

this is an introduction

INDEXING AND SLICING

In

```
import pandas as pd
```

```
min_temp_df = pd.read_csv('https://raw.githubusercontent.com/LorenzF/data-science-  
practical-approach/main/src/c2_data_preparation/data/temperatures/australia/  
melbourne/1981.csv')  
min_temp_df
```

```
      Date  Temp  
0  1981-01-01  20.7  
1  1981-01-02  17.9  
2  1981-01-03  18.8  
3  1981-01-04  14.6  
4  1981-01-05  15.8  
..      ...   ...  
360 1981-12-27  15.5  
361 1981-12-28  13.3  
362 1981-12-29  15.6  
363 1981-12-30  15.2  
364 1981-12-31  17.4
```

```
[365 rows x 2 columns]
```

```
min_temp_df.Date = pd.to_datetime(min_temp_df.Date)
```

```
min_temp_df = min_temp_df.set_index('Date')
```

```
min_temp_df.loc['1981-06-01':'1981-06-30']
```

```
      Date  Temp  
1981-06-01  11.6  
1981-06-02  10.6  
1981-06-03   9.8  
1981-06-04  11.2  
1981-06-05   5.7  
1981-06-06   7.1  
1981-06-07   2.5  
1981-06-08   3.5  
1981-06-09   4.6
```

(continues on next page)

(continued from previous page)

```
1981-06-10 11.0
1981-06-11 5.7
1981-06-12 7.7
1981-06-13 10.4
1981-06-14 11.4
1981-06-15 9.2
1981-06-16 6.1
1981-06-17 2.7
1981-06-18 4.3
1981-06-19 6.3
1981-06-20 3.8
1981-06-21 4.4
1981-06-22 7.1
1981-06-23 4.8
1981-06-24 5.8
1981-06-25 6.2
1981-06-26 7.3
1981-06-27 9.2
1981-06-28 10.2
1981-06-29 9.5
1981-06-30 9.5
```

```
min_temp_df.loc['1989-06-01':'1989-06-30'].mean()
```

```
Temp      NaN
dtype: float64
```

```
min_temp_df.resample('MS').mean()
```

```

              Temp
Date
1981-01-01  17.712903
1981-02-01  17.678571
1981-03-01  13.500000
1981-04-01  12.356667
1981-05-01   9.490323
1981-06-01   7.306667
1981-07-01   7.577419
1981-08-01   7.238710
1981-09-01  10.143333
1981-10-01  10.087097
1981-11-01  11.890000
1981-12-01  13.680645
```

```
import seaborn as sns
```

```
tip_df = sns.load_dataset('tips')
tip_df.head()
```

```

total_bill  tip    sex smoker  day    time  size
0      16.99  1.01  Female   No   Sun  Dinner     2
1      10.34  1.66   Male   No   Sun  Dinner     3
2      21.01  3.50   Male   No   Sun  Dinner     3
3      23.68  3.31   Male   No   Sun  Dinner     2
```

(continues on next page)

(continued from previous page)

```
4      24.59  3.61  Female    No  Sun  Dinner    4
```

```
tip_index_df = tip_df.set_index('day')
```

```
tip_index_df.loc['Sun']
```

```

      total_bill  tip    sex smoker   time  size
day
Sun      16.99  1.01  Female    No  Dinner    2
Sun      10.34  1.66    Male    No  Dinner    3
Sun      21.01  3.50    Male    No  Dinner    3
Sun      23.68  3.31    Male    No  Dinner    2
Sun      24.59  3.61  Female    No  Dinner    4
..          ...   ...     ...   ...   ...   ...
Sun      20.90  3.50  Female   Yes  Dinner    3
Sun      30.46  2.00    Male   Yes  Dinner    5
Sun      18.15  3.50  Female   Yes  Dinner    3
Sun      23.10  4.00    Male   Yes  Dinner    3
Sun      15.69  1.50    Male   Yes  Dinner    2

```

```
[76 rows x 6 columns]
```

```
tip_index_df = tip_df.set_index(['day', 'time'])
```

```
tip_index_df.loc[('Thur', 'Lunch')].tip.mean()
```

```

/tmp/ipykernel_25625/2537502835.py:1: PerformanceWarning: indexing past lexsort depth
may impact performance.
  tip_index_df.loc[('Thur', 'Lunch')].tip.mean()

```

```
2.767704918032786
```

```
pd.pivot_table(tip_df, values='total_bill', index='day', columns='time', aggfunc=
↳ 'median')
```

```

time  Lunch  Dinner
day
Thur   16.00  18.780
Fri    13.42  18.665
Sat      NaN  18.240
Sun      NaN  19.630

```

```
tip_df.set_index(['sex', 'time', 'smoker']).loc[('Male', 'Dinner', 'Yes')]['tip'].mean()
```

```

/tmp/ipykernel_25625/3467525553.py:1: PerformanceWarning: indexing past lexsort depth
may impact performance.
  tip_df.set_index(['sex', 'time', 'smoker']).loc[('Male', 'Dinner', 'Yes')]['tip'].
  mean()

```

```
3.123191489361702
```


Part IV

4. Data Exploration

DATA EXPLORATION

this is an introduction

Part V

5. Data Visualisation

DATA VISUALISATION

this is an introduction

Part VI

6. Machine Learning

MACHINE LEARNING

this is an introduction